

# Assignment 5

Student Name: Krishna Vamsi Koppula

Student ID: 700745021

GitHub Link:

[https://github.com/Krishnavamsikoppula/Assignment\\_5\\_NNDL\\_Summer](https://github.com/Krishnavamsikoppula/Assignment_5_NNDL_Summer)

Video Link: <https://drive.google.com/file/d/1fqzUcDYr-ZBZj8u58L92sU82E1DITz3J/view?usp=sharing>

1. Save the model and use the saved model to predict on new text data (ex, “A lot of good things are happening. We are respected again throughout the world, and that's a great thing.@realDonaldTrump”)

```
✓ 3s ▶ import pandas as pd #Basic packages for creating dataframes and loading dataset
import numpy as np

import matplotlib.pyplot as plt #Package for visualization

import re #importing package for Regular expression operations

from sklearn.model_selection import train_test_split #Package for splitting the data

from sklearn.preprocessing import LabelEncoder #Package for conversion of categorical to Numerical

from keras.preprocessing.text import Tokenizer #Tokenization
from tensorflow.keras.preprocessing.sequence import pad_sequences #Add zeros or crop based on the length
from keras.models import Sequential #Sequential Neural Network
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D #For layers in Neural Network
from keras.utils.np_utils import to_categorical
```

✓  
0s

```
# Load the dataset as a Pandas DataFrame
dataset = pd.read_csv('/content/sample_data/Sentiment.csv')

# Select only the necessary columns 'text' and 'sentiment'
mask = dataset.columns.isin(['text', 'sentiment'])
data = dataset.loc[:, mask]

# Preprocess the text data
data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))
data['text'] = data['text'].apply(lambda x: x.replace('rt', ' ')) # Remove 'rt' (Retweets)
```

```
<ipython-input-5-3b7761336ca1>:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
data['text'] = data['text'].apply(lambda x: x.lower())
<ipython-input-5-3b7761336ca1>:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))
<ipython-input-5-3b7761336ca1>:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
data['text'] = data['text'].apply(lambda x: x.replace('rt', ' ')) # Remove 'rt' (Retweets)
```

✓  
0s

```
# Define the function to create the LSTM model
def createmodel():
    model = Sequential()
    model.add(Embedding(max_features, embed_dim, input_length=X.shape[1]))
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(3, activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
    return model

# Tokenization
max_features = 2000
tokenizer = Tokenizer(num_words=max_features, split=' ')
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values)
X = pad_sequences(X)

# Label Encoding
label_encoder = LabelEncoder()
integer_encoded = label_encoder.fit_transform(data['sentiment'])
y = to_categorical(integer_encoded)
```

✓  
0s

```
[7] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

✓  
5m

```
# LSTM Model Architecture
embed_dim = 128
lstm_out = 196

model = Sequential()
model.add(Embedding(max_features, embed_dim, input_length=X.shape[1]))
model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(3, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# Model Summary
print(model.summary())

# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test), verbose=2)
```

✓  
5m

Model: "sequential"

```
Layer (type)                 Output Shape              Param #
=====
embedding (Embedding)        (None, 28, 128)          256000

lstm (LSTM)                   (None, 196)              254800

dense (Dense)                 (None, 3)                591
=====
Total params: 511,391
Trainable params: 511,391
Non-trainable params: 0

None
Epoch 1/10
291/291 - 33s - loss: 0.8253 - accuracy: 0.6465 - val_loss: 0.7528 - val_accuracy: 0.6730 - 33s/epoch - 115ms/step
Epoch 2/10
291/291 - 30s - loss: 0.6818 - accuracy: 0.7117 - val_loss: 0.7303 - val_accuracy: 0.6885 - 30s/epoch - 103ms/step
Epoch 3/10
291/291 - 30s - loss: 0.6166 - accuracy: 0.7450 - val_loss: 0.7560 - val_accuracy: 0.6765 - 30s/epoch - 103ms/step
Epoch 4/10
291/291 - 30s - loss: 0.5661 - accuracy: 0.7686 - val_loss: 0.7918 - val_accuracy: 0.6793 - 30s/epoch - 104ms/step
Epoch 5/10
291/291 - 33s - loss: 0.5184 - accuracy: 0.7875 - val_loss: 0.8440 - val_accuracy: 0.6728 - 33s/epoch - 113ms/step
Epoch 6/10
291/291 - 30s - loss: 0.4824 - accuracy: 0.8072 - val_loss: 0.8733 - val_accuracy: 0.6656 - 30s/epoch - 102ms/step
Epoch 7/10
291/291 - 31s - loss: 0.4429 - accuracy: 0.8181 - val_loss: 0.9958 - val_accuracy: 0.6496 - 31s/epoch - 107ms/step
Epoch 8/10
291/291 - 31s - loss: 0.4122 - accuracy: 0.8305 - val_loss: 1.0711 - val_accuracy: 0.6385 - 31s/epoch - 105ms/step
Epoch 9/10
291/291 - 30s - loss: 0.3754 - accuracy: 0.8472 - val_loss: 1.1761 - val_accuracy: 0.6529 - 30s/epoch - 105ms/step
Epoch 10/10
291/291 - 32s - loss: 0.3510 - accuracy: 0.8589 - val_loss: 1.1748 - val_accuracy: 0.6518 - 32s/epoch - 109ms/step
```

✓  
3s



# Evaluate the model on test data

```
score, accuracy = model.evaluate(X_test, y_test, verbose=2, batch_size=32)
print("Test Loss:", score)
print("Test Accuracy:", accuracy)
```

# Plot training and validation accuracy over epochs

```
plt.plot(history.history['accuracy'], label='Training accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation accuracy')
plt.show()
```

# Plot training and validation loss over epochs

```
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss')
plt.show()
```

✓  
3s





144/144 - 2s - loss: 1.1748 - accuracy: 0.6518 - 2s/epoch - 17ms/step

Test Loss: 1.1747829914093018



Test Accuracy: 0.6518130302429199





```
0s [1] # Save the trained model
      model.save('sentimentAnalysis.h5')

0s [11] from keras.models import load_model

      model = load_model('sentimentAnalysis.h5')

0s [12] # Define the text data to predict sentiment
      sentence = ['A lot of good things are happening. We are respected again throughout the world, and that is a great thing. @realDonaldTrump']

      # Tokenize and pad the sentence
      sentence = tokenizer.texts_to_sequences(sentence)
      sentence = pad_sequences(sentence, maxlen=28, dtype='int32', value=0)
```

✓  
0s

```
[▶] # Make predictions using the loaded model
sentiment_probs = model.predict(sentence, batch_size=1, verbose=2)[0]

# Convert sentiment probabilities to sentiment label
sentiment = np.argmax(sentiment_probs)

# Print the sentiment label
if sentiment == 0:
    print("Neutral")
elif sentiment < 0:
    print("Negative")
elif sentiment > 0:
    print("Positive")
else:
    print("Cannot be determined")
```

☞ 1/1 - 0s - 255ms/epoch - 255ms/step  
Positive

## 2. Apply GridSearchCV on the source code provided in the class

```
✓ 0s [15] from keras.wrappers.scikit_learn import KerasClassifier #importing Keras classifier
      from sklearn.model_selection import GridSearchCV #importing Grid search CV

✓ 39m [15] # Now you can proceed with the GridSearchCV
      model = KerasClassifier(build_fn=createmodel, verbose=2)
      batch_size = [10, 20, 40]
      epochs = [1, 2]
      param_grid = {'batch_size': batch_size, 'epochs': epochs}
      grid = GridSearchCV(estimator=model, param_grid=param_grid)
      grid_result = grid.fit(X_train, y_train)

      # Print the best score and best hyperparameters found by GridSearchCV
      print("Best Score: %f using %s" % (grid_result.best_score_, grid_result.best_params_))

372/372 - 35s - loss: 0.6810 - accuracy: 0.7084 - 35s/epoch - 95ms/step
93/93 - 1s - loss: 0.7344 - accuracy: 0.6724 - 1s/epoch - 15ms/step
Epoch 1/2
372/372 - 37s - loss: 0.8332 - accuracy: 0.6379 - 37s/epoch - 100ms/step
Epoch 2/2
372/372 - 34s - loss: 0.6788 - accuracy: 0.7121 - 34s/epoch - 91ms/step
93/93 - 2s - loss: 0.7464 - accuracy: 0.6864 - 2s/epoch - 21ms/step
Epoch 1/2
372/372 - 39s - loss: 0.8298 - accuracy: 0.6377 - 39s/epoch - 106ms/step
Epoch 2/2
372/372 - 37s - loss: 0.6805 - accuracy: 0.7083 - 37s/epoch - 99ms/step
93/93 - 2s - loss: 0.7453 - accuracy: 0.6868 - 2s/epoch - 17ms/step
Epoch 1/2
372/372 - 38s - loss: 0.8261 - accuracy: 0.6461 - 38s/epoch - 101ms/step
Epoch 2/2
372/372 - 35s - loss: 0.6710 - accuracy: 0.7173 - 35s/epoch - 95ms/step
93/93 - 2s - loss: 0.7940 - accuracy: 0.6555 - 2s/epoch - 19ms/step
186/186 - 23s - loss: 0.8496 - accuracy: 0.6333 - 23s/epoch - 124ms/step
47/47 - 2s - loss: 0.7789 - accuracy: 0.6547 - 2s/epoch - 33ms/step
186/186 - 25s - loss: 0.8466 - accuracy: 0.6410 - 25s/epoch - 133ms/step
```



✓  
39m



```
47/47 - 1s - loss: 0.7817 - accuracy: 0.6719 - 1s/epoch - 22ms/step
186/186 - 25s - loss: 0.8453 - accuracy: 0.6294 - 25s/epoch - 134ms/step
47/47 - 1s - loss: 0.8109 - accuracy: 0.6681 - 1s/epoch - 23ms/step
186/186 - 24s - loss: 0.8447 - accuracy: 0.6334 - 24s/epoch - 128ms/step
47/47 - 1s - loss: 0.7523 - accuracy: 0.6728 - 921ms/epoch - 20ms/step
186/186 - 23s - loss: 0.8409 - accuracy: 0.6416 - 23s/epoch - 121ms/step
47/47 - 1s - loss: 0.7801 - accuracy: 0.6695 - 1s/epoch - 27ms/step
Epoch 1/2
186/186 - 25s - loss: 0.8416 - accuracy: 0.6349 - 25s/epoch - 135ms/step
Epoch 2/2
186/186 - 22s - loss: 0.6997 - accuracy: 0.7039 - 22s/epoch - 116ms/step
47/47 - 1s - loss: 0.7354 - accuracy: 0.6762 - 1s/epoch - 24ms/step
Epoch 1/2
186/186 - 26s - loss: 0.8391 - accuracy: 0.6408 - 26s/epoch - 138ms/step
Epoch 2/2
186/186 - 22s - loss: 0.6920 - accuracy: 0.7035 - 22s/epoch - 118ms/step
47/47 - 1s - loss: 0.7383 - accuracy: 0.6891 - 1s/epoch - 22ms/step
Epoch 1/2
186/186 - 24s - loss: 0.8465 - accuracy: 0.6310 - 24s/epoch - 127ms/step
Epoch 2/2
186/186 - 21s - loss: 0.6895 - accuracy: 0.7037 - 21s/epoch - 114ms/step
47/47 - 1s - loss: 0.7488 - accuracy: 0.6815 - 896ms/epoch - 19ms/step
Epoch 1/2
186/186 - 23s - loss: 0.8477 - accuracy: 0.6293 - 23s/epoch - 125ms/step
Epoch 2/2
186/186 - 20s - loss: 0.6971 - accuracy: 0.6951 - 20s/epoch - 108ms/step
47/47 - 1s - loss: 0.7409 - accuracy: 0.6900 - 890ms/epoch - 19ms/step
Epoch 1/2
186/186 - 23s - loss: 0.8356 - accuracy: 0.6381 - 23s/epoch - 126ms/step
Epoch 2/2
186/186 - 22s - loss: 0.6834 - accuracy: 0.7095 - 22s/epoch - 118ms/step
47/47 - 1s - loss: 0.7752 - accuracy: 0.6728 - 964ms/epoch - 21ms/step
Epoch 1/2
233/233 - 29s - loss: 0.8382 - accuracy: 0.6411 - 29s/epoch - 124ms/step
Epoch 2/2
233/233 - 26s - loss: 0.6885 - accuracy: 0.7064 - 26s/epoch - 113ms/step
Best Score: 0.681911 using {'batch_size': 40, 'epochs': 2}
```

✓  
0s



```
# Plot the results of GridSearchCV
mean_scores = grid_result.cv_results_['mean_test_score']
param_batch_size = grid_result.cv_results_['param_batch_size']
param_epochs = grid_result.cv_results_['param_epochs']

plt.figure(figsize=(8, 6))
for i, batch_size in enumerate(batch_size):
    plt.plot(epochs, mean_scores[i * len(epochs): (i + 1) * len(epochs)], label=f'batch_size={batch_size}')

plt.xlabel('Number of Epochs')
plt.ylabel('Mean Test Score')
plt.title('GridSearchCV Results')
plt.legend()
plt.show()
```



GridSearchCV Results

